

Comparative Analysis of Sequential Models for Mobile Network Traffic Forecasting

Introduction

This empirical study investigates one-step-ahead Internet traffic forecasting across Milan geographical squares. The research question is how sequential models compare and whether accuracy changes with traffic characteristics across areas.

Related Work

The design follows established evidence that cellular traffic contains strong periodicity, short-term autocorrelation, nonlinear peaks, and spatial heterogeneity. Ridge autoregression provides an interpretable regularized sequence baseline; random forests capture nonlinear lag interactions; and gradient boosting provides a high-capacity additive ensemble. These models are deliberately different in inductive bias and computational cost.

Dataset and Data Preparation

The downloaded Milan dataset contains 61 daily text files and 314,412,299 records across 10,000 squares. Each record is parsed as square ID, timestamp, and Internet traffic value, with variable-length activity fields handled by taking the final field. The raw files total approximately 20 GB and are streamed line by line, so peak memory is controlled by small numeric accumulators rather than a full raw table. Data are aggregated by square and 10-minute timestamp, then represented using lagged traffic, rolling statistics, and cyclical time features.

Exploratory Analysis

The three highest-total squares were 5161, 5059, 5259. For the highest-traffic square, lag-1 autocorrelation was 0.976, lag-144 autocorrelation was 0.910, and the augmented Dickey-Fuller p-value was 3.219×10^{-15} . These results indicate strong short-term dependence and a meaningful daily cycle, supporting lag and seasonal inputs. The first-two-week comparison includes the top three squares and required squares 4159 and 4556.

Methodology

For each selected area, observations before 16 December were split chronologically into a fitting period and a final seven-day validation period. The evaluation period was 16--22 December. Ridge autoregression, random forest regression, and XGBoost regression used the same one-step feature representation: lags of 10 minutes through 24 hours, rolling means and volatility, and 10-minute/day cyclical features. Candidate hyperparameters were compared using validation RMSE, then the selected model was refit on all pre-evaluation data. No future evaluation values were used as predictors.

Results and Discussion

The tables below report MAE, RMSE, MAPE, validation RMSE, training time, and prediction time for each of the three highest-traffic areas. The tuning log records every candidate configuration and its validation metrics, so parameter changes are evidence-led rather than arbitrary. Differences across areas are interpreted alongside their temporal plots: smoother, strongly periodic areas should favour regularized lag models, while abrupt peaks provide opportunities for tree ensembles but also expose their tendency to lag unseen events.

Failure Analysis and Limitations

The most difficult periods are abrupt spikes and low-traffic intervals, where percentage error can become unstable and lag-based models cannot know an event before it appears in the history. The missing 11 December source file is a data-coverage limitation. The study also aggregates activity by square and uses only Internet traffic, so it does not model cross-area dependence, outages, or exogenous events.

Conclusion and Future Work

The comparison treats accuracy, runtime, and model transparency as joint evidence rather than selecting solely by one error metric. Future work should add the missing day if available, model all 10,000 areas with scalable distributed storage, include SMS/call covariates, and evaluate recurrent or transformer sequence models with rolling-origin validation.

Selected evidence figures




Figure 1. Distribution of total Internet traffic across all geographical squares.




Figure 2. First-two-week traffic comparison for the three highest areas and squares 4159 and 4556.




Figure 3. Highest-area time series and autocorrelation diagnostics.




Figure 4. Representative one-step forecast for square 5161.

Top geographical areas

Square	Total Internet traffic
5161	12644487.277
5059	11111075.043
5259	10293534.429

Forecasting results

Area	Model	MAE	RMSE	MAPE	Train s	Predict s
5161	RidgeAR	89.554	130.243	10.47%	0.007	0.0020
5161	RandomForest	92.997	138.905	8.71%	1.287	0.0739
5161	XGBoost	96.897	146.123	9.24%	0.457	0.0056
5059	RidgeAR	78.075	105.753	9.26%	0.015	0.0057
5059	RandomForest	74.972	107.517	7.14%	1.459	0.0798
5059	XGBoost	76.636	111.523	7.65%	0.471	0.0053
5259	RidgeAR	71.886	100.738	8.36%	0.009	0.0027
5259	RandomForest	66.458	97.647	7.23%	1.481	0.0807
5259	XGBoost	67.468	96.526	7.68%	1.085	0.0109

Iterative tuning log

Area	Model	Candidate	Validation RMSE	Parameters
5161	RidgeAR	1	161.316	{'alpha': 1.0}
5161	RidgeAR	2	162.168	{'alpha': 10.0}
5161	RidgeAR	3	170.096	{'alpha': 100.0}
5161	RandomForest	1	169.221	{'n_estimators': 120, 'max_depth': 12, 'min_samples_leaf': 2}
5161	RandomForest	2	170.077	{'n_estimators': 200, 'max_depth': 18, 'min_samples_leaf': 2}
5161	XGBoost	1	181.340	{'n_estimators': 250, 'max_depth': 5, 'learning_rate': 0.05}
5161	XGBoost	2	188.317	{'n_estimators': 400, 'max_depth': 6, 'learning_rate': 0.03}
5059	RidgeAR	1	128.391	{'alpha': 1.0}
5059	RidgeAR	2	128.686	{'alpha': 10.0}

Area	Model	Candidate	Validation RMSE	Parameters
5059	RidgeAR	3	131.915	{'alpha': 100.0}
5059	RandomForest	1	152.382	{'n_estimators': 120, 'max_depth': 12, 'min_samples_leaf': 2}
5059	RandomForest	2	155.287	{'n_estimators': 200, 'max_depth': 18, 'min_samples_leaf': 2}
5059	XGBoost	1	155.100	{'n_estimators': 250, 'max_depth': 5, 'learning_rate': 0.05}
5059	XGBoost	2	162.530	{'n_estimators': 400, 'max_depth': 6, 'learning_rate': 0.03}
5259	RidgeAR	1	115.180	{'alpha': 1.0}
5259	RidgeAR	2	116.843	{'alpha': 10.0}
5259	RidgeAR	3	126.663	{'alpha': 100.0}
5259	RandomForest	1	109.229	{'n_estimators': 120, 'max_depth': 12, 'min_samples_leaf': 2}
5259	RandomForest	2	109.329	{'n_estimators': 200, 'max_depth': 18, 'min_samples_leaf': 2}
5259	XGBoost	1	108.657	{'n_estimators': 250, 'max_depth': 5, 'learning_rate': 0.05}
5259	XGBoost	2	106.705	{'n_estimators': 400, 'max_depth': 6, 'learning_rate': 0.03}

References

[1] G. Barlacchi et al., A multi-source dataset of urban life in the city of Milan and the Province of Trentino, Scientific Data, 2015. [2] Harvard Dataverse, Milan telecommunications dataset, DOI: 10.7910/DVN/EGZHFV. [3] T. Hastie, R. Tibshirani, and J. Friedman, The Elements of Statistical Learning, Springer.